Eight diagnostics, four evaluations: the model ranks everywhere, and calibrates only on its own data

Read across a row for one evaluation's whole profile; down a column to compare one diagnostic across evaluations. Dashed lines mark each scale's reference value.
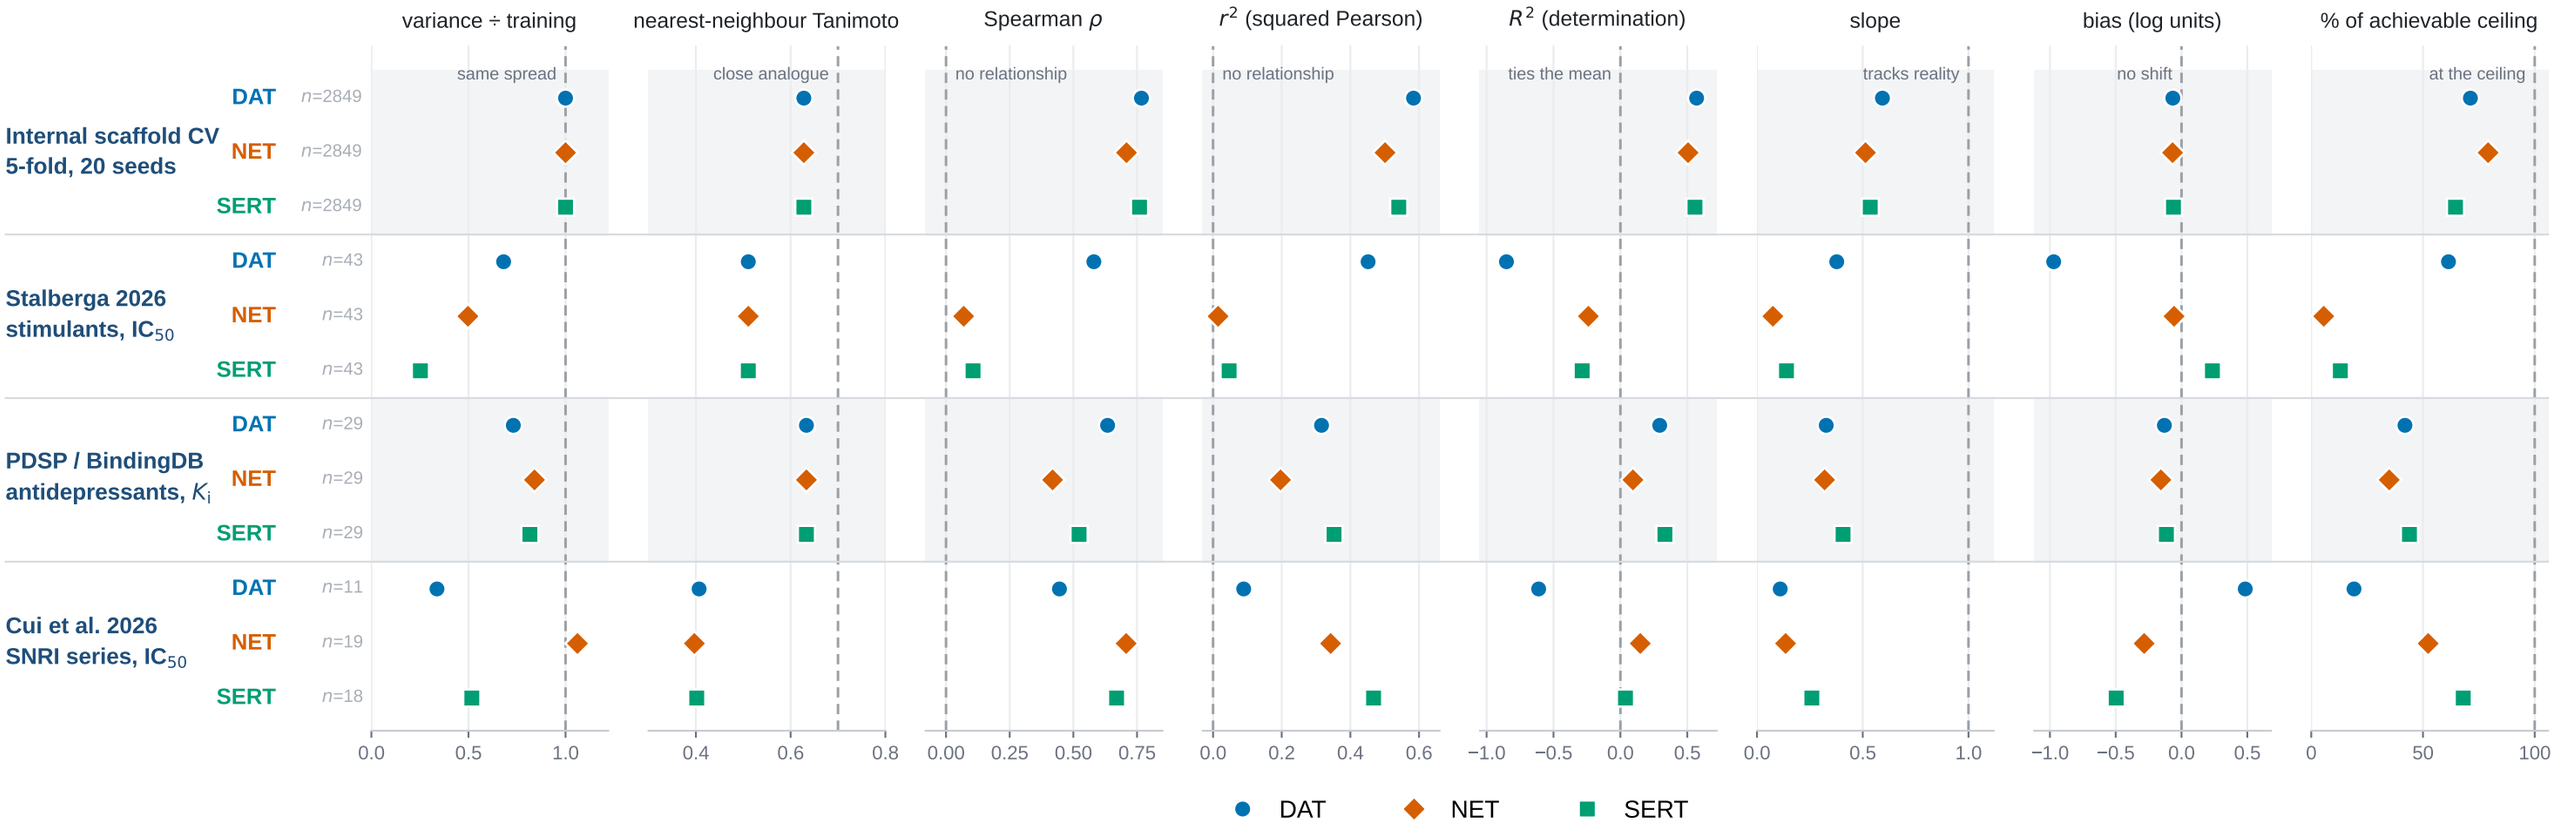
Internal r^2 and % of ceiling are the published 20-seed means; the other internal diagnostics come from the seed-0 fold realisation. The internal variance ratio is 1.00 by construction: the held-out folds are drawn from the training distribution itself.