Eight diagnostics, five evaluations: a large internal gain that does not leave the training distribution
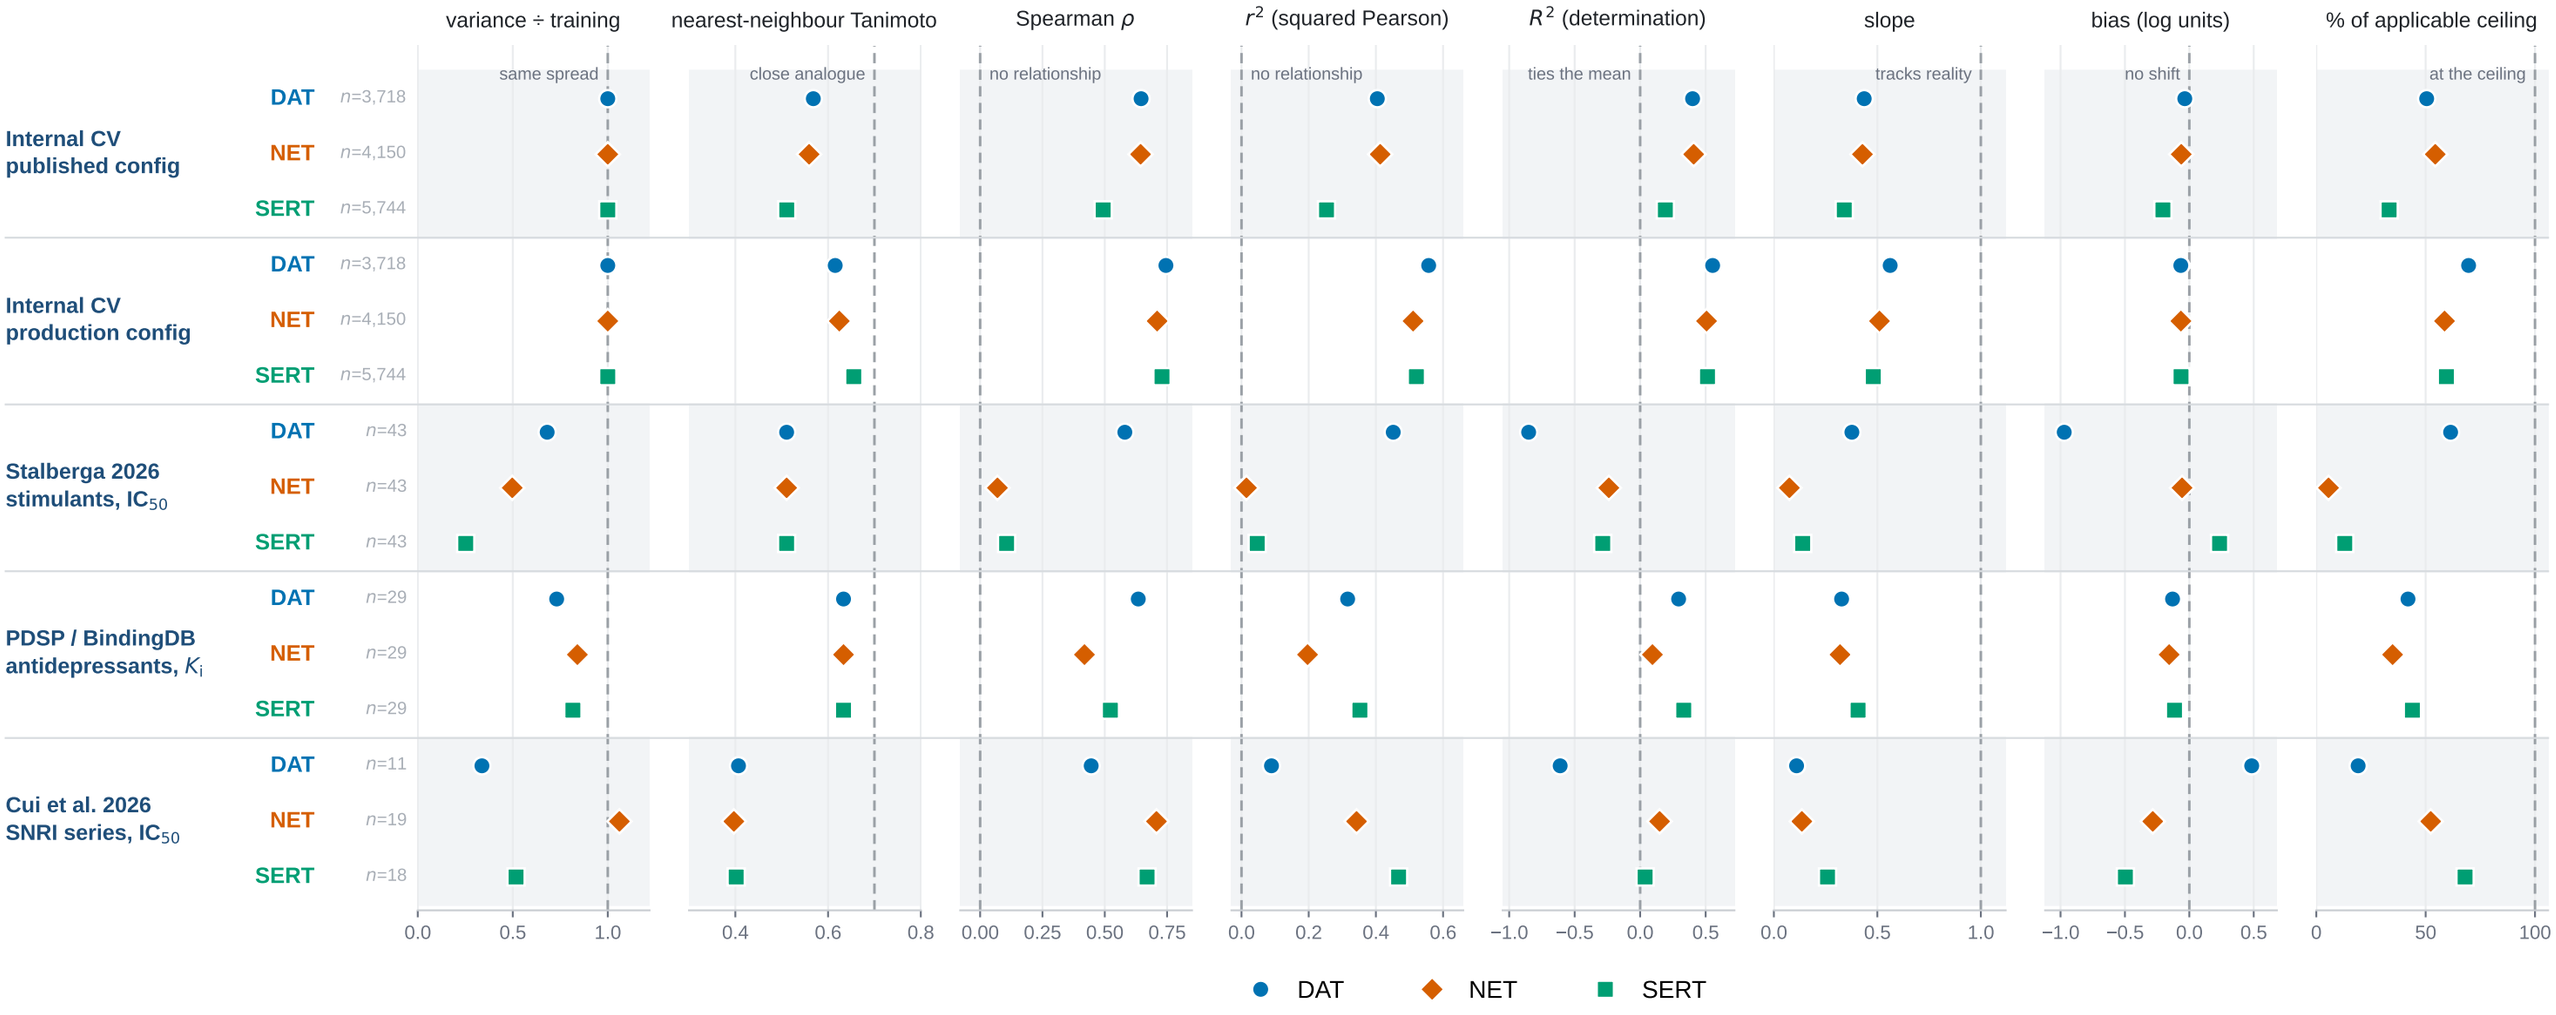
Read across a row for one evaluation's whole profile; down a column to compare one diagnostic across evaluations. Dashed lines mark each scale's reference value. *n* is the number of values scored. Both internal blocks are 20 seeds × 5 scaffold folds with every diagnostic re-measured per seed; the published config trains on the 2,849 compounds measured at all three transporters, the production config on every compound with that one target (3,367 / 3,702 / 5,338) plus a binary assay-modality feature. Internal variance ratio is 1.00 by construction, and nearest-neighbour Tanimoto is fold-wise. The last panel uses each row's own applicable ceiling: the published model predicts a measurement of unknown assay, the production model one of known assay, and each external set is scored against its own response variance.